

Batch: B2 Roll No.: 1611087

Experiment No. 03

Grade: AA / AB / BB / BC / CC / CD / DD

Title: Database Tuning

Objective: Tuning the database to improve system performance

Expected Outcome of Experiment:

CO1 : Design and tune database.

Books/ Journals/ Websites referred:

1. *Elmasri & Navathe “ fundamentals of Database Systems” V edition. PEARSON Education.*
 2. *Korth, Silberschatzsu darshan “Database systems, concepts” 5th edition McGraw Hill.*
 3. *Raghu Ramkrishnan & Johannes Gehrke “Database Management System” Tata McGraw Hill. III edition.*
-

Pre Lab/ Prior Concepts: Database, ER diagram, Relation mapping, SQL

Implementation Details:

Perform following;

1. Index tuning
2. Query Tuning
3. Database Tuning

K. J. Somaiya College of Engineering, Mumbai-77

(Autonomous College Affiliated to University of Mumbai)

And mention why tuning is required? How it is performed and what is the effect of tuning?
for all mentioned three categories of tuning.

INDEX TUNING

- Normal Select operation which accesses all records

✓ Showing rows 0 - 1 (2 total, Query took 0.0020 seconds.)

```
SELECT * FROM enquiry WHERE project_name = 'avighna'
```

☐ Profiling [Edit inline] [Edit] [Explain SQL] [Create PHP code] [Refresh]

☐ Show all | Number of rows: 25 | Filter rows: Search this table | Sort by key: None

+ Options

		user_name	user_email	project_name	enquiry_id	user_number
<input type="checkbox"/>	Edit Copy Delete	tarun lohana	tarun.lohana@somaiya.edu	avighna	1	9876543217
<input type="checkbox"/>	Edit Copy Delete	nishit patel	noshit.pateli@somaiya.edu	avighna	6	7777688881

☐ Check all | With selected: Edit Copy Delete Export

☐ Show all | Number of rows: 25 | Filter rows: Search this table | Sort by key: None

- Explain select operation thus proving that all records are accessed

K. J. Somaiya College of Engineering, Mumbai-77

(Autonomous College Affiliated to University of Mumbai)



Show query box

⚠ Current selection does not contain a unique column. Grid edit, checkbox, Edit, Copy and Delete features are not available.

Your SQL query has been executed successfully.

```
EXPLAIN SELECT * FROM `enquiry` WHERE project_name = 'avighna'
```

[\[Edit inline\]](#) [\[Edit \]](#) [\[Skip Explain SQL \]](#) [\[Analyze Explain at mariadb.org\]](#) [\[Create PHP code \]](#)

+ Options

id	select_type	table	type	possible_keys	key	key_len	ref	rows	Extra
1	SIMPLE	enquiry	ALL	NULL	NULL	NULL	NULL	6	Using where

- Created index on 'project_name' column (Primary Index)



Show query box

✓ MySQL returned an empty result set (i.e. zero rows). (Query took 0.1950 seconds.)

```
CREATE INDEX project_name_index on enquiry(project_name)
```

[\[Edit inline\]](#) [\[Edit \]](#) [\[Create PHP code \]](#)

- Indexed query (We can see that less number of records are fetched thus reducing processing time)

K. J. Somaiya College of Engineering, Mumbai-77

(Autonomous College Affiliated to University of Mumbai)



Show query box

⚠ Current selection does not contain a unique column. Grid edit, checkbox, Edit, Copy and Delete features are not available.

Your SQL query has been executed successfully.

```
EXPLAIN SELECT * FROM enquiry where project_name = 'avighna'
```

[\[Edit inline\]](#) [\[Edit \]](#) [\[Skip Explain SQL \]](#) [\[Analyze Explain at mariadb.org\]](#) [\[Create PHP code \]](#)

+ Options

id	select_type	table	type	possible_keys	key	key_len	ref	rows	Extra
1	SIMPLE	enquiry	ref	project_name_index	project_name_index	102	const	2	Using index condition

- Created index on 'user_name' column (Clustered Index)

← Server: 127.0.0.1 » Database: construction » Table: enquiry ⚙



Show query box

✓ MySQL returned an empty result set (i.e. zero rows). (Query took 0.1060 seconds.)

```
create index user_name_index on enquiry(user_name)
```

[\[Edit inline\]](#) [\[Edit \]](#) [\[Create PHP code \]](#)

- Indexed query (Same explanation as above)

K. J. Somaiya College of Engineering, Mumbai-77

(Autonomous College Affiliated to University of Mumbai)

Server: 127.0.0.1 » Database: construction » Table: enquiry

⏮ Browse ⏭ Structure Ⓞ SQL 🔍 Search ➕ Insert 📤 Export 📂 Import ▼ More

Show query box

⚠ Current selection does not contain a unique column. Grid edit, checkbox, Edit, Copy and Delete features are not available.

✓ Your SQL query has been executed successfully.

```
explain SELECT * FROM `enquiry` WHERE user_name = 'shivam sansare'
```

[Edit inline] [Edit] [Skip Explain SQL] [Analyze Explain at mariadb.org] [Create PHP code]

+ Options

id	select_type	table	type	possible_keys	key	key_len	ref	rows	Extra
1	SIMPLE	enquiry	ref	user_name_index	user_name_index	102	const	2	Using index condition

- Create index on 2 columns (Secondary Index)

Server: 127.0.0.1 » Database: construction » Table: enquiry

⏮ Browse ⏭ Structure Ⓞ SQL 🔍 Search ➕ Insert 📤 Export 📂 Import ▼ More

Show query box

✓ MySQL returned an empty result set (i.e. zero rows). (Query took 0.1340 seconds.)

```
create INDEX user_name_and_project_name on enquiry(user_name, project_name)
```

[Edit inline] [Edit] [Create PHP code]

- Explain query on 2 indexes(We can see that less rows are fetched from)

K. J. Somaiya College of Engineering, Mumbai-77

(Autonomous College Affiliated to University of Mumbai)

Browser Structure SQL Search Insert Export Import More

Show query box

⚠ Current selection does not contain a unique column. Grid edit, checkbox, Edit, Copy and Delete features are not available.

Your SQL query has been executed successfully.

```
explain SELECT * FROM `enquiry` WHERE user_name = 'akshay kotak'
```

[Edit inline] [Edit] [Skip Explain SQL] [Analyze Explain at mariadb.org] [Create PHP code]

+ Options

id	select_type	table	type	possible_keys	key	key_len	ref	rows
1	SIMPLE	enquiry	ref	user_name_index,user_name_and_project_name	user_name_index	102	const	

- Explain query on 2 indexes

Show query box

⚠ Current selection does not contain a unique column. Grid edit, checkbox, Edit, Copy and Delete features are not available.

Your SQL query has been executed successfully.

```
explain SELECT * FROM `enquiry` WHERE user_name = 'akshay kotak' and project_name = 'palava'
```

[Edit inline] [Edit] [Skip Explain SQL] [Analyze Explain at mariadb.org] [Create PHP code]

+ Options

id	select_type	table	type	possible_keys	key	key_len	ref	rows	Extra
1	SIMPLE	enquiry	ref	project_name_index,user_name_index,user_name_and_p...	user_name_index	102	const	1	Using index condition; Using where

- Explain query on 2 indexes

K. J. Somaiya College of Engineering, Mumbai-77

(Autonomous College Affiliated to University of Mumbai)

Your SQL query has been executed successfully.

```
explain SELECT * FROM `enquiry` WHERE user_name = 'shivam sansare' and project_name = 'palava'
```

[\[Edit inline\]](#) [\[Edit\]](#) [\[Skip Explain SQL\]](#) [\[Analyze Explain at mariadb.org\]](#) [\[Create PHP code\]](#)

+ Options

id	select_type	table	type	possible_keys	key	key_len	ref	rows	Extra
1	SIMPLE	enquiry	ref	project_name_index,user_name_index,user_name_and_p...	user_name_and_project_name	204	const,const	1	Using index conditi

Query results operations

Print Copy to clipboard Create view

QUERY TUNING

- SELECT (enquiry_id) FROM enquiry
WHERE project_name IN (SELECT project_name
FROM project
WHERE project_name is NOT NULL)

SQL PERFORMANCE TUNING

recommendations:

```
SELECT (enquiry_id) FROM enquiry  
WHERE project_name IN ( SELECT project_name  
FROM project  
WHERE project_name is NOT NULL )
```

ERRORS SECTION

- ORA-00942: table or view does not exist

ALSO CONSIDER ABOUT

Instead of IN you should also try to use OR or EXISTS for subqueries

K. J. Somaiya College of Engineering, Mumbai-77

(Autonomous College Affiliated to University of Mumbai)

Why tuning is necessary?

Ans: Since we are fetching records from inner query by using IN operator which doesn't check if inner query returned rows or not which isn't a good option.

How tuning is performed?

Ans: By including EXISTS clause we tune the query for optimal performance.

- `SELECT (user_id) FROM login l, documents d
WHERE user_id = architectId`

SQL PERFORMANCE TUNING

recommendations:

```
SELECT (user_id) FROM login l, documents d  
WHERE user_id = architectId
```

ERRORS SECTION

- ORA-00942: table or view does not exist

Why tuning is necessary?

Ans: No tuning is necessary for this query, it is optimal.

- `SELECT (user_id) FROM login
WHERE user_id in (select architectId FROM documents
WHERE project_name exists (select project_name
FROM project
where price_per_sqft > 10000))`

SQL PERFORMANCE TUNING

recommendations:

```
SELECT (user_id) FROM login
WHERE user_id in ( select architectId FROM documents
WHERE project_name exists ( select
project_name
FROM project
where price_per_sqft
> 10000) )
```

ERRORS SECTION

- ORA-00920: invalid relational operator

ALSO CONSIDER ABOUT

Instead of IN you should also try to use OR or EXISTS for subqueries
Instead of EXISTS you should also try to use IN

- select designation
from login L
where L.user_id in (select purpose
from accounts A
)

K. J. Somaiya College of Engineering, Mumbai-77

(Autonomous College Affiliated to University of Mumbai)

```
select designation
from login L
where L.user_id in ( select purpose
from accounts A
)
```

ERRORS SECTION

- ORA-00942: table or view does not exist

ALSO CONSIDER ABOUT

Instead of IN you should also try to use OR or EXISTS for subqueries

Why tuning is necessary?

Ans: Instead of comparing each records from inner query with outer query, one can use JOIN condition and do the same with optimal performance.

How tuning is performed?

Ans: By using JOIN relational operator.

- SELECT document_name
FROM documents d, project p
WHERE a.project_name = p.project_name;

SQL PERFORMANCE TUNING

recommendations:

```
SELECT document_name  
FROM documents d, project p  
WHERE a.project_name = p.project_name
```

ERRORS SECTION

- ORA-00942: table or view does not exist

No tuning required. Query is already optimised.

- SELECT user_name, remaining_flats
FROM enquiry e, project p
WHERE e.project_name = p.project_name;

K. J. Somaiya College of Engineering, Mumbai-77

(Autonomous College Affiliated to University of Mumbai)

```
SELECT user_name, remaining_flats
FROM enquiry e, project p
WHERE e.project_name = p.project_name
```

ERRORS SECTION

- ORA-00942: table or view does not exist

-
- ```
SELECT user_name, count(1)
FROM enquiry e
WHERE e.project_name in (select project_name
 from project)
```

## SQL PERFORMANCE TUNING recommendations:

```
SELECT user_name,count(1)
FROM enquiry e
WHERE e.project_name in (select project_name
from project)
```

---

### ERRORS SECTION

---

- ORA-00942: table or view does not exist

---

### ALSO CONSIDER ABOUT

Instead of IN you should also try to use OR or EXISTS for subqueries

Why tuning is necessary?

Ans: Since count(\*) counts all the rows of the result table for every column which is redundant. We can do so by specifying count(1) too.

How tuning is performed?

By using count(1) instead of count(\*) we reduced the query processing time.

## DATABASE TUNING

Database tuning means tuning the database for optimal performance while querying the database.

Like making less use of number to string conversions etc.

Here,

## NO DENORMALIZATION IS REQUIRED

**Conclusion:**

**Hence the various tuning like index,query,database tuning been successfully understood and implemented**

**Post Lab Descriptive Questions:**

1. What are the factors that influence Physical Database Design ?

**Ans:**

Physical design is an activity where the goal is not only to create the appropriate structuring of data in storage, but also to do so in a way that guarantees good performance. For a given conceptual schema, there are many physical design alternatives in a given DBMS. It is not possible to make meaningful physical design decisions and performance analyses until the database designer knows the mix of queries, transactions, and applications that are expected to run on the database. This is called the job mix for the particular set of database system applications. The database administrators/designers must analyze these applications, their expected frequencies of invocation, any timing constraints on their execution speed, the expected frequency of update operations, and any unique constraints on attributes. We discuss each of these factors next.

**Design Criteria**

- Response Time  
Elapsed Time between submitting a database transaction for execution and receiving a response
- Space Utilization  
Storage space used by database files and their access path structures
- Transaction throughput
  - Average number of transactions/minute
  - Must be measured under peak conditions

## **K. J. Somaiya College of Engineering, Mumbai-77**

(Autonomous College Affiliated to University of Mumbai)

### **A. Analyzing the Database Queries and Transactions.**

Before undertaking the physical database design, we must have a good idea of the intended use of the database by defining in a high-level form the queries and transactions that are expected to run on the database. For each retrieval query, the following information about the query would be needed:

- The files that will be accessed by the query.
- The attributes on which any selection conditions for the query are specified.
- Whether the selection condition is an equality, inequality, or a range condition.
- The attributes on which any join conditions or conditions to link multiple tables or objects for the query are specified.
- The attributes whose values will be retrieved by the query.

For each update operation or update transaction, the following information would be needed:

- The files that will be updated.
- The type of operation on each file (insert, update, or delete).
- The attributes on which selection conditions for a delete or update are specified.
- The attributes whose values will be changed by an update operation.

Again, the attributes listed in item 3 are candidates for access structures on the files, because they would be used to locate the records that will be updated or deleted. On the other hand, the attributes listed in item 4 are candidates for *avoiding an access structure*, since modifying them will require updating the access structures.

### **B. Analyzing the Expected Frequency of Invocation of Queries and Transactions.**

Besides identifying the characteristics of expected retrieval queries and update transactions, we must consider their expected rates of invocation. This frequency information, along with the attribute information collected on each query and transaction, is used to compile a cumulative list of the expected frequency of use for all queries and transactions. This is expressed as the expected frequency of using each attribute in each file as a selection attribute or a join attribute, over all the queries and transactions. Generally, for large volumes of processing, the informal *80–20 rule* can be used: approximately 80 percent of the processing is accounted for by only 20 percent of the queries and transactions. Therefore, in practical situations, it is rarely



necessary to collect exhaustive statistics and invocation rates on all the queries and transactions; it is sufficient to determine the 20 percent or so most important ones.

### **C.Analyzing the Time Constraints of Queries and Transactions.**

Some queries and transactions may have stringent performance constraints. For example, a transaction may have the constraint that it should terminate within 5 seconds on 95 percent of the occasions when it is invoked, and that it should never take more than 20 seconds. Such timing constraints place further priorities on the attributes that are candidates for access paths. The selection attributes used by queries and transactions with time constraints become higher-priority candidates for primary access structures for the files, because the primary access structures are generally the most efficient for locating records in a file.

### **D.Analyzing the Expected Frequencies of Update Operations.**

A minimum number of access paths should be specified for a file that is frequently updated, because updating the access paths themselves slows down the update operations. For example, if a file that has frequent record insertions has 10 indexes on 10 different attributes, each of these indexes must be updated whenever a new record is inserted. The overhead for updating 10 indexes can slow down the insert operations.

### **E.Analyzing the Uniqueness Constraints on Attributes.**

Access paths should be specified on all *candidate key* attributes—or sets of attributes—that are either the primary key of a file or unique attributes. The existence of an index (or other access path) makes it sufficient to only search the index when checking this uniqueness constraint, since all values of the attribute will exist in the leaf nodes of the index. For example, when inserting a new record, if a key attribute value of the new record *already exists in the index*, the insertion of the new record should be rejected, since it would violate the uniqueness constraint on the attribute.

**K. J. Somaiya College of Engineering, Mumbai-77**

**(Autonomous College Affiliated to University of Mumbai)**

**Date:** \_\_\_\_\_



